

# Linked Lists: Locking, Lock-Free, and Beyond ...



BROWN

Maurice Herlihy

CS176

Fall 2003

# Linked Lists

- We can make effective spin locks
  - Correct
  - Work well under contention
- Are we done with concurrent data structures?

# Not Over Yet

- Contention
  - Solved by MCS or CLH locks
- Sequential Bottleneck
  - No "bag of tricks"
- Linked Lists
  - Simple data structure
  - Good testbed

# Set Interface

- Unordered collection of objects
- No duplicates
- Methods
  - Add a new object
  - Remove an object
  - Test if object is present

# List-Based Sets

```
public interface Set {  
    public boolean add(Object x);  
    public boolean remove(Object x);  
    public boolean contains(Object x);  
}
```

# List Entry

```
public class Entry {  
    public Object object;  
    public int key;  
    public Entry next;  
}
```

# List Entry

```
public class Entry {  
    public Object object;  
    public int key;  
    public Entry next;  
}
```

Object of interest

# List Entry

```
public class Entry {  
    public Object object;  
    public int key;  
    public Entry next;  
}
```

Sort by key value  
(usually hash code)



# List Entry

```
public class Entry {  
    public Object object;  
    public int key;  
    public Entry next;  
}
```

Sorting makes it  
easy to detect absence




# List Entry

```
public class Entry {  
    public Object object;  
    public int key;  
    public Entry next;  
}
```


Reference to next entry

# List-Based Set




**Sentinel node never deleted  
(minimum possible key)**

# Adding an Entry



# Removing an Entry



# Nu, What About Concurrency?

- Our bag of tricks
  - Coarse-grained locks
  - Fine-grained locks
  - Optimistic synchronization
  - Lock-free synchronization

# Coarse-Grained Locking

- Easy, same as synchronized methods
  - "One lock to rule them all ..."
- Simple, clearly correct
  - Deserves respect!
- Works poorly with contention
  - Queue locks help
  - But bottleneck still an issue

# Fine-grained Locking

- Requires careful thought
  - "Do not meddle in the affairs of wizards, for they are subtle and quick to anger"
- Split object into pieces
  - Each piece has own lock
  - Methods that work on disjoint pieces need not exclude each other




# Optimistic Synchronization

- Requires very careful thought
  - "Do not meddle in the affairs of dragons, for you are crunchy and taste good with ketchup."
- Try it without synchronization
  - If you win, you win
  - If not, try it again with synchronization


# Lock-Free Synchronization

- Dump locking altogether ...
  - "You take the red pill and you stay in Wonderland and I show you how deep the rabbit-hole goes"
- No locks, just native atomic methods
  - Usually `compareAndSet`


# Hand-over-Hand locking




# Removing an Entry



# Removing an Entry



# Uh-oh



# Problem


- To delete entry  $b$ 
  - Swing entry  $a$ 's next field to  $c$
- Problem is,
  - Someone could delete  $c$  concurrently

# Insight


- If an entry is locked
  - No one can delete entry's *successor*
- If a thread locks
  - Entry to be deleted
  - And its predecessor
  - Then it works




# Hand-Over-Hand Again



# Removing an Entry



# Removing an Entry



# Remove method

```
public boolean remove(Object object) {  
    int key = object.hashCode();  
    Entry predEntry, currEntry;  
    try {  
        ...  
    } finally {  
        currEntry.unlock();  
        predEntry.unlock();  
    }  
}
```

# Remove method

```
public boolean remove(Object object) {  
    int key = object.hashCode();  
    Entry prevEntry, currEntry;  
    try {  
        ...  
    } finally {  
        currEntry.unlock();  
        prevEntry.unlock();  
    }  
}
```

**Key used to order entry**



# Remove method

```
public boolean remove(Object object) {  
    int key = object.hashCode();  
    Entry predEntry, currEntry;  
    try {  
        ...  
    } finally {  
        currEntry.unlock();  
        prevEntry.unlock();  
    }  
}
```

**Predecessor and current entries**



# Remove method

```
public boolean remove(Object object) {  
    int key = object.hashCode();  
    Entry prevEntry, currEntry;  
    try {  
        ...  
    } finally {  
        currEntry.unlock();  
        predEntry.unlock();  
    }  
}
```

**Make sure  
locks released**

# Remove method

```
public boolean remove(Object object) {  
    int key = object.hashCode();  
    Entry prevEntry, currEntry;  
    try {  
        ...  
    } finally {  
        currEntry.unlock();  
        prevEntry.unlock();  
    }  
}
```

**Everything else**






# Remove method

```
try {  
    predEntry = this.head;  
    predEntry.lock();  
    currEntry = predEntry.next;  
    currEntry.lock();  
    ...  
} finally { ... }
```

# Remove method

```
try {  
    predEntry = this.head;  
    predEntry.lock();  
    currEntry = predEntry.next;  
    currEntry.lock();  
    ...  
} finally { ... }
```


lock previous



# Remove method

```
try {  
    prevEntry = this.head;  
    prevEntry.lock();  
    currEntry = prevEntry.next;  
    currEntry.lock();  
    ...  
} finally { ... }
```


Lock current



# Remove method


```
try {  
    prevEntry = this.head;  
    prevEntry.lock();  
    currEntry = prevEntry.next;  
    currEntry.lock();  
    ...  
} finally { ... }
```

Traversing list



# Remove: searching


```
while (currEntry.key <= key) {  
    if (object == currEntry.object) {  
        predEntry.next = currEntry.next;  
        return true;  
    }  
    predEntry.unlock();  
    predEntry = currEntry;  
    currEntry = currEntry.next;  
    currEntry.lock();  
}  
return false;
```



# Remove: searching

```
while (currEntry.key <= key) {  
    if (object == currEntry.object) {  
        prevEntry.next = currEntry.next;  
        return true;  
    }  
    prevEntry.unlock();  
    prevEntry = currEntry;  
    currEntry = currEntry.next;  
    currEntry.lock();  
}  
return false;
```


**Search key range**



# Remove: searching

```
while (currEntry.key <= key) {  
    if (object == currEntry.object) {  
        prevEntry.next = currEntry.next;  
        return true;  
    }  
    prevEntry = currEntry;  
    currEntry = currEntry.next;  
}  
return false;
```


**At start of each loop:  
currEntry and predEntry  
locked**



# Remove: searching

```
while (currEntry.key <= key) {  
  if (object == currEntry.object) {  
    predEntry.next = currEntry.next;  
    return true;  
  }  
  prevEntry.unlock();  
  prevEntry = currEntry;  
  currEntry = currEntry.next;  
  currEntry.lock();  
}
```

**If entry found, remove it**






# Remove: searching

```
while (currEntry.key <= key) {  
  if (object == currEntry.object) {  
    predEntry.next = currEntry.next;  
    return true;  
  }  
  prevEntry.unlock();  
  prevEntry = currEntry;  
  currEntry = currEntry.next;  
  currEntry.lock();  
}
```


**If entry found, remove it**



# Remove: searching

Unlock predecessor

```
while (currEntry.key <= key) {  
    if (object == currEntry.object) {  
        prevEntry.next = currEntry.next;  
        return true;  
    }  
    predEntry.unlock();  
    prevEntry = currEntry;  
    currEntry = currEntry.next;  
    currEntry.lock();  
}  
return false;
```




# Remove: searching

**Only one entry locked!**

```
while (currEntry.key <= key) {  
    if (object == currEntry.object) {  
        prevEntry.next = currEntry.next;  
        return true;  
    }  
    prevEntry = currEntry;  
    currEntry = currEntry.next;  
    currEntry.lock();  
}  
return false;
```


**prevEntry.unlock();**



# Remove: searching


```
... remove(Object o) {  
    ... (object = currEntry.object) {  
        prevEntry.next = currEntry.next;  
        return true;  
    }  
    prevEntry.unlock();  
    predEntry = currEntry;  
    currEntry = currEntry.next;  
    currEntry.lock();  
}  
return false;
```

**demote current**



# Remove: searching


```
while (currEntry.key <= key) {  
    Find and lock new current (ct) {  
        prevEntry.next = currEntry.next;  
        return true;  
    }  
    prevEntry.unlock();  
    prevEntry = currEntry;  
    currEntry = currEntry.next;  
    currEntry.lock();  
}  
return false;
```



# Remove: searching

**Lock invariant restored**

```
if (object == currEntry.object) {  
    prevEntry.next = currEntry.next;  
    return true;  
}  
prevEntry.unlock();  
prevEntry = currEntry;  
currEntry = currEntry.next;  
currEntry.lock();  
}  
return false;
```



# Remove: searching

```
while (currEntry.key <= key) {  
    if (object == currEntry.object) {  
        prevEntry.next = currEntry.next;  
        return true;  
    }  
    prevEntry.unlock();  
    prevEntry = currEntry;  
    currEntry = currEntry.next;  
    currEntry.lock();  
}
```

**Otherwise, error 404, dude!**

**return false;**



# Why does this work?

- To remove entry  $e$ 
  - Must lock  $e$
  - Must lock  $e$ 's predecessor
- Therefore, if you lock an entry
  - It can't be removed
  - And neither can its successor



# Lock and Load

- To move to successor entry for  $e$ 
  - Lock  $e$
  - Lock  $e.next$  Until next entry identified and locked
  - Unlock  $e$  Don't release  $e$
- While traversing
  - $e$  cannot be removed
  - $e.next$  cannot be removed

# Adding Entries

- To add entry  $e$ 
  - Must lock predecessor
  - Must lock successor
- Neither can be deleted
  - (Is successor lock actually required?)


# Drawbacks

- Better than coarse-grained lock
  - Threads can traverse in parallel
- Still not ideal
  - Long chain of acquire/release
  - Inefficient


# Optimistic Synchronization

- Find entries without locking
- Lock entries
- Check that everything is OK


# Removing an Entry




# What Can Go Wrong?



# Check that Entry is Still Accessible




# What Can Go Wrong?






# What Can Go Wrong?




# Check that Entries Still Adjacent



# Correctness

- If
  - Entry  $b$  and Entry  $c$  both locked
  - Entry  $b$  still accessible
  - Entry  $c$  still successor to  $b$
- Then
  - Neither will be deleted
  - OK to delete and return `true`

# Removing an Absent Entry



# Correctness

- If
  - Entry  $b$  and Entry  $d$  both locked
  - Entry  $b$  still accessible
  - Entry  $d$  still successor to  $b$
- Then
  - Neither will be deleted
  - No thread can add  $c$  after  $b$
  - OK to return `false`


# Validation

```
private boolean
  validate(Entry predEntry,
           Entry currEntry) {
  Entry entry = head;
  while (entry.key <= predEntry.key) {
    if (entry == predEntry)
      return predEntry.next == currEntry;
    entry = entry.next;
  }
  return false;
}
```

# Validation


```
private boolean  
validate(Entry predEntry,  
         Entry currEntry) {  
    Entry entry = head;  
    while (entry.key <= predEntry.key) {  
        if (entry == predEntry) {  
            return predEntry.next == null;  
        }  
        entry = entry.next;  
    }  
    return false;  
}
```

**Predecessor &  
current entries**



# Validation

```
private boolean  
validate(Entry predEntry,  
         Entry currEntry) {  
    Entry entry = head;  
    while (entry.key <= predEntry.key) {  
        if (entry == predEntry)  
            return predEntry.next == currEntry;  
        entry = entry.next;  
    }  
    return false;  
}
```




**Start at the  
beginning**





# Validation

```
private boolean
validate(Entry predEntry,
        Entry currEntry) {
    Entry entry = head;
    while (entry.key <= predEntry.key) {
        if (entry == predEntry)
            return predEntry.next == currEntry;
        entry = entry.next;
    }
    return false;
}
```




**Search range of keys**



# Validation

```
private boolean
validate(Entry predEntry,
        Entry currEntry) {
    Entry entry = head;
    while (entry.key <= predEntry.key) {
        if (entry == predEntry)
            return predEntry.next == currEntry;
        entry = entry.next;
    }
    return false;
}
```




**Predecessor reachable**



# Validation

```
private boolean
validate(Entry predEntry,
        Entry currEntry) {
    Entry entry = head;
    while (entry.key <= predEntry.key) {
        if (entry == predEntry)
            return predEntry.next == currEntry;
        entry = entry.next;
    }
    return false;
}
```




**Is current entry next?**

# Validation

```
private boolean  
validate(Entry predEntry,  
         Entry currEntry) {  
    Entry entry = head;  
    while (entry.key <= predEntry.key) {  
        if (entry == predEntry)  
            return predEntry.next == currEntry;  
        entry = entry.next;  
    }  
    return false;  
}
```


Otherwise move on



# Validation

```
private boolean validate(Entry predEntry,
                        Entry currEntry) {
    Entry entry = head;
    while (entry.key <= predEntry.key) {
        if (entry == predEntry)
            return predEntry.next == currEntry;
        entry = entry.next;
    }
    return false;
}
```

**Predecessor not reachable**




# Remove: searching

```
public boolean remove(Object object) {  
    int key = object.hashCode();  
    retry: while (true) {  
        Entry predEntry = this.head;  
        Entry currEntry = predEntry.next;  
        while (currEntry.key <= key) {  
            if (object == currEntry.object)  
                break;  
            predEntry = currEntry;  
            currEntry = currEntry.next;  
        } ...  
    }
```

# Remove: searching

```
public boolean remove(Object object) {  
    int key = object.hashCode();  
    retry: while (true) {  
        Entry prevEntry = this.head;  
        Entry currEntry = prevEntry.next;  
        while (currEntry.key <= key) {  
            if (object == currEntry.object) {  
                break;  
            }  
            prevEntry = currEntry;  
            currEntry = currEntry.next;  
        }  
        ...  
    }  
}
```




**Search key**



# Remove: searching

```
public boolean remove(Object object) {  
    int key = object.hashCode();  
    retry: while (true) {  
        Entry prevEntry = this.head;  
        Entry currEntry = prevEntry;  
        while (currEntry.key <= key)  
            if (object == currEntry.obj)  
                break;  
        prevEntry = currEntry;  
        currEntry = currEntry.next;  
    } ...
```


**Retry on synchronization conflict**





# Remove: searching

```
public boolean remove(Object object) {  
    int key = object.hashCode();  
    retry: while (true) {  
        Entry predEntry = this.head;  
        Entry currEntry = predEntry.next;  
        while (currEntry.key <= key)  
            if (object == currEntry.obj)  
                break;  
        prevEntry = currEntry;  
        currEntry = currEntry.next;  
    }  
}
```




**Examine predecessor and current entries**

# Remove: searching

```
public boolean remove(Object object) {  
    int key = object.hashCode();  
    retry: while (true) {  
        Entry prevEntry = this.head;  
        Entry currEntry = prevEntry.next;  
        while (currEntry.key <= key) {  
            if (object == currEntry.obj)  
                break;  
            prevEntry = currEntry;  
            currEntry = currEntry.next;  
        }  
        ...  
    }  
}
```


**Search by key**



# Remove: searching

```
public boolean remove(Object object) {  
    int key = object.hashCode();  
    retry: while (true) {  
        Entry prevEntry = this.head;  
        Entry currEntry = prevEntry.next;  
        while (currEntry.key <= key) {  
            if (object == currEntry.object) break;  
            prevEntry = currEntry;  
            currEntry = currEntry.next;  
        }  
    }  
}
```


**Stop if we find object**



# Remove: searching

```
public void remove(Object object) {  
    int hashCode = object.hashCode();  
    retry: while (true) {  
        Entry prevEntry = this.head;  
        Entry currEntry = prevEntry.next;  
        while (currEntry.key <= hashCode) {  
            if (object == currEntry.object) {  
                break;  
            }  
            predEntry = currEntry;  
            currEntry = currEntry.next;  
        }  
        ...  
    }  
}
```

**Move along**



# On Exit from Loop

- If object is present
  - currEntry holds object
  - predEntry just before currEntry
- If object is absent
  - currEntry has first higher key
  - predEntry just before currEntry
- Assuming no synchronization problems

# Remove Method

```
try {  
    predEntry.lock(); currEntry.lock();  
    if (validate(predEntry, currEntry) {  
        if (currEntry.object == object) {  
            predEntry.next = currEntry.next;  
            return true;  
        } else {  
            return false;  
        }  
    } finally {  
        predEntry.unlock();  
        currEntry.unlock();  
    }  
}
```



# Remove Method

```
try {
```

```
    predEntry.lock(); currEntry.lock();  
    if (validate(predEntry, currEntry) {  
        if (currEntry.object == object) {  
            predEntry.next = currEntry.next;  
            return true;  
        } else {  
            return false;  
        }  
    }  
}
```

```
    } finally {  
        predEntry.unlock();  
        currEntry.unlock();  
    }  
}
```

**Always unlock**



# Remove Method

```
try {  
    predEntry.lock(); currEntry.lock();  
    if (validate(predEntry, currEntry) {  
        if (currEntry.object == object) {  
            predEntry.next = currEntry.next;  
            return true;  
        } else {  
            return false;  
        }  
    }  
    finally {  
        predEntry.unlock();  
        currEntry.unlock();  
    }  
}
```

**Lock both entries**





# Remove Method

```
try {  
    predEntry.lock(); currEntry.lock();  
    if (validate(predEntry, currEntry)) {  
        if (currEntry.object == object) {  
            predEntry.next = currEntry.next;  
            return true;  
        } else {  
            return false;  
        }  
    }  
    finally {  
        predEntry.unlock();  
        currEntry.unlock();  
    }  
}
```

**Check for synchronization  
conflicts**



# Remove Method

```
try {  
    predEntry.lock(); currEntry.lock();  
    if (validate(predEntry, currEntry) {  
        if (currEntry.object == object) {  
            predEntry.next = currEntry.next;  
            return true;  
        } else {  
            return false;  
        }  
    }  
    finally {  
        predEntry.unlock();  
        currEntry.unlock();  
    }  
}
```

**Object found,  
remove entry**



# Remove Method

```
try {
    predEntry.lock(); currEntry.lock();
    if (validate(predEntry, currEntry) {
        if (currEntry.object == object) {
            predEntry.next = currEntry.next;
            return true;
        } else {
            return false;
        }
    }
} finally {
    predEntry.unlock();
    currEntry.unlock();
}
```

**Object not found**



# So Far, So Good

- Much less lock acquisition/release
  - Performance
  - Concurrency
- Problems
  - Need to traverse list twice
  - contains() method acquires locks
    - Most common method call


# Marked Lists

- Remove Method
  - Scans list (as before)
  - Locks predecessor & current (as before)
  - Marks current entry as removed (new!)
  - Redirects predecessor's next (as before)


# Marked Lists

- All Methods
  - Scan list
  - Do not scan past marked entry
  - Instead, start over from list head

# Business as Usual




# Interference






# Interference




# Interference



# Marking a Node

- AtomicMarkableReference **class**
  - Java.util.concurrent.atomic **package**



# Extracting Information

```
Public Object get(boolean[]);
```

# Extracting Information

```
Public object get(boolean[]);
```

Returns  
reference

Stores mark  
at index 0



# Extracting Information

```
public boolean isMarked();
```

**Value of  
mark**



# Changing State

```
Public boolean compareAndSet(  
    Object expectedRef,  
    Object updateRef,  
    boolean expectedMark,  
    boolean updateMark);
```

# Changing State

If this is the current  
reference ...

```
Public boolean compareAndSet(  
    Object expectedRef,  
    Object updateRef,  
    boolean expectedMark,  
    boolean updateMark);
```

And this is the  
current mark ...





# Changing State

...then change to this  
new reference ...

```
Public boolean compareAndSet(  
    Object expectedRef,  
    Object updateRef,  
    boolean expectedMark,  
    boolean updateMark);
```

... and this new  
mark



# Changing State

```
public boolean attemptMark(  
    Object expectedRef,  
    boolean updateMark);
```

# Changing State

```
public boolean attemptMark(  
    Object expectedRef,  
    boolean updateMark);
```

**If this is the current  
reference ...**



# Changing State

```
public boolean attemptMark(  
    Object expectedRef,  
    boolean updateMark);
```

.. then change to  
this new mark.



# Aside

- We will see that it is often useful to tag pointers with
  - Boolean values
  - Integer values
- Sometimes to mark, and sometimes to ensure pointers are unique

# List Validate Method

```
private boolean  
    validate(Entry predEntry,  
             Entry currEntry) {  
    return  
        (!predEntry.next.isMarked()) &&  
        (!currEntry.next.isMarked()) &&  
        (predEntry.next.getReference()  
         == currEntry);  
}
```

# List Validate Method

```
private boolean Predecessor not removed  
    validate(Entry predEntry,  
             Entry currEntry) {  
    return  
        (!predEntry.next.isMarked()) &&  
        (!currEntry.next.isMarked()) &&  
        (predEntry.next.getReference()  
         == currEntry);  
}
```

# List Validate Method

```
private boolean Current not removed  
    validate(Entry predEntry,  
             Entry currEntry) {  
    return  
        (!predEntry.next.isMarked()) &&  
        (!currEntry.next.isMarked()) &&  
        (predEntry.next.getReference()  
         == currEntry);  
}
```



# List Validate Method

```
private boolean Next field unchanged  
    validate(Entry predEntry,  
             Entry currEntry) {  
    return  
        (!predEntry.next.isMarked()) &&  
        (!currEntry.next.isMarked()) &&  
        (predEntry.next.getReference()  
         == currEntry);  
}
```

# Remove: searching

```
public boolean remove(Object object) {  
    ...  
    while (currEntry.key <= key) {  
        Entry nextEntry =  
            (Entry)currEntry.next.get(mark);  
        if (mark[0])  
            continue retry;  
        if (object == currEntry.object)  
            break;  
        predEntry = currEntry;  
        currEntry = currEntry.next;  
    } ...  
}
```

# Remove: searching

```
public boolean remove(Object object) {  
    ...  
    while (currEntry.key <= key) {  
        Entry nextEntry =  
            (Entry)currEntry.next.get(mark);  
        if (mark[0])  
            continue retry;  
        if (object == currEntry.object)  
            break;  
        predEntry = c  
        currEntry = currEntry.next;  
    } ...  
}
```

**Atomically read reference  
& mark**



# Remove: searching

**Panic if entry removed**

```
... act object) {  
...  
while (currEntry.key <= key) {  
    Entry nextEntry =  
        (Entry)currEntry.next.get(mark);  
    if (mark[0])  
        continue retry;  
    if (object == currEntry.object)  
        break;  
    predEntry = currEntry;  
    currEntry = currEntry.next;  
} ...
```



# Evaluation

- Good:
  - Contains method doesn't need to lock
  - Uncontended calls don't re-traverse
- Bad
  - Contended calls do re-traverse
  - Traffic jam if one thread delays

# Traffic Jam

- Any concurrent data structure based on mutual exclusion has a weakness
- If one thread
  - Enters critical section
  - And "eats the big muffin" (stops running)
    - Cache miss, page fault, descheduled ...
    - Software error, ...
  - Everyone else using that lock is stuck!

# Lock-Free Data Structures



- No matter what ...
  - Some thread will complete method call
  - Even if others halt at malicious times
- Implies that
  - You can't use locks (why?)
  - Um, that's why they call it lock-free

# Lock-Free $\neq$ Wait-Free


- Wait-free synchronization
  - Every method call eventually finishes
  - What everyone really wants
- Lock-free synchronization
  - Some method call eventually finishes
  - What we are usually willing to pay for
    - Starvation rare in practice ...




# Lock-Free Lists

- Next logical step
- Eliminate locking entirely
- Use only `compareAndSet()`
- Invented by Maged Michael, 2003


# Adding an Entry




# Removing an Entry



# Removing an Entry



# Look Familiar?




# Problem

- Method updates entry's next field
- After entry has been removed

# Solution


- Use AtomicMarkableReference
- Remove in two steps
  - Set mark bit in next field
  - Redirect predecessor's pointer
- CAS
  - Fails if mark bit set (entry removed)

# Removing an Entry







# Removing an Entry



# Removing an Entry




# Removing an Entry



# Traversing the List

- Q: what do you do when you find a "logically" deleted entry in your path?
- A: finish the job.
  - CAS the predecessor's next field
  - Proceed (repeat as needed)

# Lock-Free Traversal




# The Find Method

```
pred, curr, next = find(object);
```

# The Find Method

```
pred, curr, succ = find(object);
```

At some instant,




or ...

# The Find Method

```
pred, curr, succ = find(object);
```

At some instant,

**object** not in list





# Remove

```
public boolean remove(Object object) {  
    while (true) {  
        pred, curr, succ = find(object);  
        if (curr == null)  
            return false;  
        if (!curr.next.attemptMark(succ,  
                                   true))  
            continue;  
        pred.next.compareAndSet(curr, succ,  
                                 false, false);  
        return true;  
    }  
}
```

# Remove

```
public boolean remove(Object object) {  
    while (true) {  
        pred, curr, succ = find(object);  
        if (curr == null)  
            return false;  
        if (!curr.next.attemptMark(succ,  
                                   true))  
            continue;  
        pred.next.compareAndSet(curr, succ,  
                                  false, false);  
        return true;  
    }  
}
```

**Keep trying**



# Remove

```
public boolean remove(Object object) {  
    while (true) {  
        pred, curr, succ = find(object);  
        if (curr == null)  
            return false;  
        if (!curr.next.attemptMark(succ,  
                                   true))  
            continue;  
        pred.next.compareAndSet(curr, succ,  
                                   false, false);  
        return true;  
    }  
}
```

**Find alleged neighbors**



# Remove

```
public boolean remove(Object object) {  
    while (true) {  
        pred, curr, succ = find(object);  
        if (curr == null)  
            return false;  
        if (!curr.next.attemptMark(succ,  
                                   true))  
            continue;  
        pred.next.compareAndSet(curr, succ,  
                                  false, false);  
        return true;  
    }  
}
```

**She's not there ...**



# Remove

```
Try to mark entry as deleted
while (true) {
    pred, curr, succ = find(object);
    if (curr == null)
        return false;
    if (!curr.next.attemptMark(succ,
                              true))
        continue;
    pred.next.compareAndSet(curr, succ,
                           false, false);
    return true;
}
```



# Remove

**If it doesn't work, just retry**


```
...ect) {  
while (true) {  
    pred, curr, succ = find(object);  
    if (curr == null)  
        return false;  
    if (!curr.next.attemptMark(succ,  
                                true))  
        continue;  
    pred.next.compareAndSet(curr, succ,  
                             false, false);  
    return true;  
}}
```



# Remove

If it works, our job is (essentially) done


```
while (curr != null) {  
    pred, curr, succ = find(object);  
    if (curr == null)  
        return false;  
    if (!curr.next.attemptMark(succ,  
                                true))  
        continue;  
    pred.next.compareAndSwap(curr, succ);  
    return true;  
}}
```



# Remove

Try to advance reference  
(if we don't succeed, someone else did).

```
if (curr == null)  
    return false;  
if (!curr.next.compareAndSet(curr, succ,  
    true))  
    continue;
```



```
pred.next.compareAndSet(curr, succ,  
    false, false);  
return true;
```



# Add

```
public boolean add(Object object) {  
    while (true) {  
        pred, curr, succ = find(object);  
        if (curr != null)  
            return false;  
        Entry entry = new Entry(object);  
        entry.next = new AMR(succ, false);  
        if (pred.next.CAS(succ, entry,  
                           false, false))  
            return true;  
    }  
}
```

# Add

```
public boolean add(Object object) {  
    while (true) {  
        pred, curr, succ = find(object);  
        if (curr != null)  
            return false;  
        Entry entry = new Entry(object);  
        entry.next = new AMR(succ, false);  
        if (pred.next.CAS(succ, entry,  
                           false, false))  
            return true;  
    }  
}
```

**Object already there.**




# Add

```
public boolean add(Object  
while (true) {  
    pred, curr, succ= find(  
    if (curr != null)  
        return false;
```

```
    Entry entry = new Entry(object);  
    entry.next = new AMR(succ, false);
```


```
    if (pred.next.CAS(succ, entry,  
                      false, false))
```

**create new entry**



# Add

```
public boolean add(Object object) {
```



```
    find(object);
```

**Install new entry**

```
    Entry entry = new Entry(object);  
    entry.next = new AMR(succ, false);
```

```
    if (pred.next.CAS(succ, entry,  
                      false, false))
```

```
        return true;
```



```
}}
```

# Contains

```
public boolean contains(Object obj){  
    while (true) {  
        prev, curr, succ = find(object);  
        return (curr != null);  
    }  
}
```

# Contains

```
public boolean contains(Object obj){  
    while (true) {  
        prev, curr, succ = find(object);  
        return (curr != null);  
    }  
}
```



Did we find anything?

# Find

```
private Entry, Entry, Entry  
    find(Object object) {  
    Entry pred, curr, succ;  
    boolean[] pmark = new boolean[1];  
    boolean[] cmark = new boolean[1];  
    int key = object.hashCode();  
    tryAgain: while (true) {  
        ...  
    }  
}
```

# Find

The entries we seek

```
private Entry, Entry, Entry  
find(Object object) {  
    Entry pred, curr, succ;  
    boolean[] pmark = new boolean[1];  
    boolean[] cmark = new boolean[1];  
    int key = object.hashCode();  
    tryAgain: while (true) {  
        ...  
    }  
}
```





# Find

```
private Entry, Entry, Entry  
    find(Object object) {  
    Entry pred, curr, succ;  
    boolean[] pmark = new boolean[1];  
    boolean[] cmark = new boolean[1];  
    int key = object.hashCode(),  
    tryAgain: while (true) {  
        ...  
    }  
}
```

Deleted bits for pred  
and curr



# Find

```
private Entry, Entry, Entry  
    find(Object object) {  
    Entry pred, curr, succ;  
    boolean[] pmark = new boolean[1];  
    boolean[] cmark = new boolean[1];  
    int key = object.hashCode();  
    tryAgain: while (true) {  
        ...  
    }  
}
```

If list changes while  
traversed, start over

# Find

```
private Entry, Entry, Entry  
    find(Object object) {  
    Entry pred, curr, succ;  
    boolean[] pmark = new boolean[1];  
    boolean[] cmark = new boolean[1];  
    int key = object.hashCode();  
    tryAgain: while (true) {  
        ...  
    }  
}
```


**Lock-Free because we start over  
only if someone else makes progress**



# Find

```
tryAgain: while (true) {  
    pred = this.head.getReference();  
    curr = pred.next.get(pmark);  
    while (true) {  
        ...  
    }  
}
```

**Start with first two entries**



# Find


```
tryAgain: while (true) {  
    pred = this.head.getReference();  
    curr = pred.next.get(pmark);  
    while (true) {  
        ...  
    }  
}
```

**Move down the list**

# Find

...

```
while (true) {  
  if (curr == null)  
    return pred, null, succ;  
  succ = curr.next.get(cmark);  
  int ckey = curr.key;  
  if (isChanged(pred.next))  
    continue tryAgain;  
}}
```



**Fell off the end of the list**

# Find

```
...  
while (true) {  
    if (curr == null)  
        return pred, null, succ;  
    succ = curr.next.get(cmark);  
    int ckey = curr.key;  
    if (isChanged(pred.next))  
        continue tryAgain;  
}}}
```

**Get ref to successor and  
current deleted bit**


# Find

Panic if predecessor's next  
field changed

```
...  
while (true) {  
    if (curr == null)  
        return pred, null, succ;  
    succ = curr.next.get(cmark);  
    int ckey = curr.key;  
    if (isChanged(pred.next))  
        continue tryAgain;  
}}}
```



# Find




```
while (true) {  
    ...  
    if (!cmark[0]) {  
        if (curr.object == object)  
            return pred, curr, succ;  
        else if (ckey <= key) {  
            pred = curr;  
        } else  
            return prev, null, curr;  
    } else {  
        ...  
    }  
}
```

**If current node is not deleted**


# Find

```
while (true) {  
    ...  
    if (!cmark[0]) {  
        if (curr.object == object)  
            return pred, curr, succ;  
        else if (ckey <= key) {  
            pred = curr;  
        } else  
            return prev, null, curr;  
    } else {  
        ...  
    }  
}
```



**Object found**

# Find



```
while (true) {  
    ...  
    if (!cmark[0]) {  
        if (curr.object == object)  
            return pred, curr, succ;  
        else if (ckey <= key) {  
            pred = curr;  
        } else  
            return prev, null, curr;  
    } else {  
        ...  
    }  
}
```


**Keep looking**

# Find

```
while (true) {  
    ...  
    if (!cmark[0]) {  
        if (curr.object == object)  
            return pred, curr, succ;  
        else if (ckey <= key) {  
            pred = curr;  
        } else  
            return prev, null, curr;  
    } else {  
        ...  
    }  
}
```

**Not there, give up**

# Find



```
...  
while (true) {  
    ...  
    if (!cmark[0]) {
```

```
        ...  
        } else {
```


```
            if (pred.next.compareAndSet(  
                curr, succ, false, false))  
                continue;  
            else  
                continue tryAgain;
```

```
}
```

**Current entry is  
logically deleted**



# Find




Try to redirect predecessor's  
next reference

```
while (true) {  
    ...  
    if (!cmark[0]) {  
        ...  
    } else {  
        if (pred.next.compareAndSet(  
            curr, succ, false, false))  
            continue;  
        else  
            continue tryAgain;  
    }  
}
```



# Find



On success, keep going,  
on failure, start over




```
if (pred.next.compareAndSet(  
    curr, succ, false, false))  
    continue;  
else  
    continue tryAgain;
```

# Summary


- Coarse-grained locking
- Fine-grained locking
- Optimistic synchronization
- Lock-free synchronization




# Scratch




# Scratch




# Scratch




# Scratch



# Scratch



# Scratch



# Removing an Entry

